






OnePlus 15, Llama-3.2-1B, 9737-token prompt, 1024 output tokens, GPU, K=1024. One real request per pull, cooled and pinned. The seed happened to draw no exploration step after the warm start, so the non-greedy arms rest on one pull per context.