




(a) GPU: cache budget vs energy per token




(b) GPU: cache budget vs throughput, every cell




(c) GPU: clock cap vs energy per token



(d) where a request's energy goes



(e) v1 controller, same command line, battery state varied



(f) answer quality per tier, Llama-1B

