




(a) energy against GPU clock cap



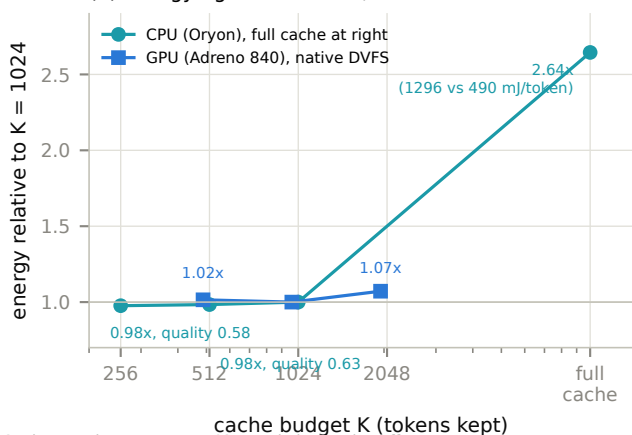
(b) energy against CPU clock: flat within 7%



(c) energy against output length: linear in tokens



(d) energy against cache K, relative to K = 1024



OnePlus 15, Llama-3.2-1B, 9737-token prompt. (a) whole-request cells, n=3 at 4096 tokens and n=5 to 11 at 1024, cooled, charging off.
(b) decode-phase rail power over tok/s, CPU clock campaign. (c) cost-table lines against cells from the controller, scheduler and bandit proofs.
(d) CPU: six-generation soak (full cache) and controller proof (K tiers, 1024 tokens); GPU: cache axis, n=3 per K, clock not pinned, 4096 tokens.