


throughput energy efficiency cache reduction

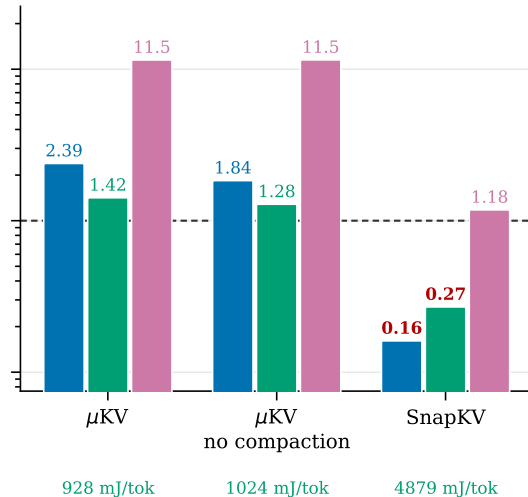
× better than the full cache

Llama-3.2-1B



full cache: 308 mJ/tok

Phi-3-mini



full cache: 1315 mJ/tok

All three axes normalised so >1 is better than the full cache. μ KV clears 1.0 on every axis; SnapKV falls below it on two of three — it pays the eviction cost and collects neither the speed nor the energy. No prior KV-eviction paper reports the energy axis at all.