




(a) energy: prediction vs meter



(b) time: prefill (grey) + decode



(c) energy per output token



arms, as battery state told to the scheduler and the plan it chose: control = mains, GPU 1200 MHz, 4096 tokens. 80% = GPU 1200 MHz, 4096 tokens. 40% = GPU cap 902 MHz, output cap 1024. 15% = GPU cap 902 MHz, output cap 512. 15% 4096 fix = told 15%, caller fixed the length: clock cap only. The CPU is chosen only when the GPU is unavailable (see the decision map).