


INFERENCE PASS


llama.cpp, CPU or Adreno GPU, one request

→ data - - - → control ·····→ heat - ···→ measure



CONTROL PLANE

THERMAL 2 Hz



ENERGY per request, 0.3 s

