

INFERENCE PASS llama.cpp, CPU or Adreno GPU, one request

