



Same policy, two backends: KV/W decides what eviction is worth (Llama-3.2-1B, ctx 16384, each policy at its own published budget)

Phone CPU — attention is compute-bound:
eviction pays

Phone GPU — weights dominate the read:
eviction ties



μ KV (attention-driven, seq-level) positional, seq-level (own budget) per-head evictors [†] old build, unpinned, timing-only