


On-device LLM workload signatures on OnePlus 15 (Phi-3-mini, full-cache) — power, memory, and temperature vs. time


(A) Long-context retrieval

8K prompt → 83-tok answer · prefill-bound



(B) Interactive assistant

4K prompt → 100-tok reply · short turn



(C) Sustained generation

7.5K prompt → 2048-tok output · decode-bound

